

```
!pip install pandas numpy matplotlib seaborn
```

Requirement already satisfied: pandas in /usr/local/lib/python3.11/dist-packages (2.2.2)
Requirement already satisfied: numpy in /usr/local/lib/python3.11/dist-packages (2.0.2)
Requirement already satisfied: matplotlib in /usr/local/lib/python3.11/dist-packages (3.10.0)
Requirement already satisfied: seaborn in /usr/local/lib/python3.11/dist-packages (0.13.2)
Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/python3.11/dist-packages (from pandas) (2.8.2)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.11/dist-packages (from pandas) (2025.1)
Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.11/dist-packages (from pandas) (2025.1)
Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.11/dist-packages (from matplotlib) (1.3.1)
Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.11/dist-packages (from matplotlib) (0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.11/dist-packages (from matplotlib) (4.56.0)
Requirement already satisfied: kiwisolver>=1.3.1 in /usr/local/lib/python3.11/dist-packages (from matplotlib) (1.4.8)
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.11/dist-packages (from matplotlib) (24.2)
Requirement already satisfied: pillow>=8 in /usr/local/lib/python3.11/dist-packages (from matplotlib) (11.1.0)
Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.11/dist-packages (from matplotlib) (3.2.1)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.11/dist-packages (from python-dateutil>=2.8.2->pandas) (1.17.0)

```
import pandas as pd
```

```
# Load the dataset with the correct file name
df = pd.read_csv('greendestination.csv')
```

```
# Display the first few rows
print("First 5 rows of the dataset:")
print(df.head())
```

```
# Check data types and missing values
print("\nDataset Info:")
print(df.info())
```

```
# Basic statistics
print("\nBasic Statistics:")
print(df.describe())
```

| | | | | | |
|-----|-----------|-------------|-----------|----------|-----|
| std | 9.135373 | 403.509100 | 8.106864 | 1.024165 | 0.0 |
| min | 18.000000 | 102.000000 | 1.000000 | 1.000000 | 1.0 |
| 25% | 30.000000 | 465.000000 | 2.000000 | 2.000000 | 1.0 |
| 50% | 36.000000 | 802.000000 | 7.000000 | 3.000000 | 1.0 |
| 75% | 43.000000 | 1157.000000 | 14.000000 | 4.000000 | 1.0 |
| max | 60.000000 | 1499.000000 | 29.000000 | 5.000000 | 1.0 |

| | EmployeeNumber | EnvironmentSatisfaction | HourlyRate | JobInvolvement | \ |
|-------|----------------|-------------------------|-------------|----------------|---|
| count | 1470.000000 | 1470.000000 | 1470.000000 | 1470.000000 | |
| mean | 1024.865306 | 2.721769 | 65.891156 | 2.729932 | |
| std | 602.024335 | 1.093082 | 20.329428 | 0.711561 | |
| min | 1.000000 | 1.000000 | 30.000000 | 1.000000 | |
| 25% | 491.250000 | 2.000000 | 48.000000 | 2.000000 | |
| 50% | 1020.500000 | 3.000000 | 66.000000 | 3.000000 | |
| 75% | 1555.750000 | 4.000000 | 83.750000 | 3.000000 | |
| max | 2068.000000 | 4.000000 | 100.000000 | 4.000000 | |

| | JobLevel | ... | RelationshipSatisfaction | StandardHours | \ |
|-------|-------------|-----|--------------------------|---------------|---|
| count | 1470.000000 | ... | 1470.000000 | 1470.0 | |
| mean | 2.063946 | ... | 2.712245 | 80.0 | |
| std | 1.106940 | ... | 1.081209 | 0.0 | |
| min | 1.000000 | ... | 1.000000 | 80.0 | |
| 25% | 1.000000 | ... | 2.000000 | 80.0 | |
| 50% | 2.000000 | ... | 3.000000 | 80.0 | |
| 75% | 3.000000 | ... | 4.000000 | 80.0 | |
| max | 5.000000 | ... | 4.000000 | 80.0 | |

| | StockOptionLevel | TotalWorkingYears | TrainingTimesLastYear | \ |
|-------|------------------|-------------------|-----------------------|---|
| count | 1470.000000 | 1470.000000 | 1470.000000 | |
| mean | 0.793878 | 11.279592 | 2.799320 | |
| std | 0.852077 | 7.780782 | 1.289271 | |
| min | 0.000000 | 0.000000 | 0.000000 | |
| 25% | 0.000000 | 6.000000 | 2.000000 | |
| 50% | 1.000000 | 10.000000 | 3.000000 | |
| 75% | 1.000000 | 15.000000 | 3.000000 | |
| max | 3.000000 | 40.000000 | 6.000000 | |

| | WorkLifeBalance | YearsAtCompany | YearsInCurrentRole | \ |
|-------|-----------------|----------------|--------------------|---|
| count | 1470.000000 | 1470.000000 | 1470.000000 | |
| mean | 2.761224 | 7.008163 | 4.229252 | |
| std | 0.706476 | 6.126525 | 3.623137 | |
| min | 1.000000 | 0.000000 | 0.000000 | |
| 25% | 2.000000 | 3.000000 | 2.000000 | |
| 50% | 3.000000 | 5.000000 | 3.000000 | |
| 75% | 3.000000 | 9.000000 | 7.000000 | |
| max | 4.000000 | 40.000000 | 18.000000 | |

| | YearsSinceLastPromotion | YearsWithCurrManager |
|-------|-------------------------|----------------------|
| count | 1470.000000 | 1470.000000 |
| mean | 2.187755 | 4.123129 |
| std | 3.222430 | 3.568136 |
| min | 0.000000 | 0.000000 |
| 25% | 0.000000 | 2.000000 |
| 50% | 1.000000 | 3.000000 |
| 75% | 3.000000 | 7.000000 |
| max | 15.000000 | 17.000000 |

[8 rows x 26 columns]

```
# Check for duplicates
print("Number of duplicate rows:", df.duplicated().sum())
```


```
# Drop unnecessary columns (e.g., EmployeeCount, Over18, StandardHours are constant)
df = df.drop(columns=['EmployeeCount', 'Over18', 'StandardHours'])
```

```
# Convert categorical variables to appropriate types
df['Attrition'] = df['Attrition'].map({'Yes': 1, 'No': 0}) # Convert to binary (1 for Yes, 0 for No)
```

```
# Number of duplicate rows: 0
```

```
# Calculate attrition rate
total_employees = len(df)
employees_left = df['Attrition'].sum()
attrition_rate = (employees_left / total_employees) * 100
```

```
print(f"Total Employees: {total_employees}")
print(f"Employees Who Left: {employees_left}")
print(f"Attrition Rate: {attrition_rate:.2f}%")
```

 Total Employees: 1470
Employees Who Left: 237
Attrition Rate: 16.12%

```
# Step 5: Analyze Factors Influencing Attrition (Improved)
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report
from sklearn.preprocessing import StandardScaler
```

```
# Correlation matrix for numeric variables (unchanged)
numeric_cols = ['Age', 'YearsAtCompany', 'MonthlyIncome', 'Attrition']
correlation = df[numeric_cols].corr()
print("\nCorrelation Matrix:")
print(correlation)
```

```
# Visualize correlations (unchanged)
plt.figure(figsize=(8, 6))
sns.heatmap(correlation, annot=True, cmap='coolwarm', fmt='.2f')
plt.title('Correlation Matrix')
plt.show()
```

```
# Boxplots to visualize distributions (unchanged)
plt.figure(figsize=(15, 5))
plt.subplot(1, 3, 1)
sns.boxplot(x='Attrition', y='Age', data=df)
plt.title('Age vs Attrition')
```

```
plt.subplot(1, 3, 2)
sns.boxplot(x='Attrition', y='YearsAtCompany', data=df)
plt.title('YearsAtCompany vs Attrition')
```

```
plt.subplot(1, 3, 3)
sns.boxplot(x='Attrition', y='MonthlyIncome', data=df)
plt.title('MonthlyIncome vs Attrition')
plt.show()
```

```
# Logistic Regression with scaling and class weighting
X = df[['Age', 'YearsAtCompany', 'MonthlyIncome']]
y = df['Attrition']
```

```
# Scale the features
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
```

```
# Split the data
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.3, random_state=42)
```

```
# Train the model with class weighting
model = LogisticRegression(class_weight='balanced', random_state=42)
model.fit(X_train, y_train)
```

```
# Print coefficients
print("\nLogistic Regression Coefficients (after scaling):")
for feature, coef in zip(X.columns, model.coef_[0]):
    print(f"{feature}: {coef:.4f}")
```



```
# Model performance
y_pred = model.predict(X_test)
print("\nClassification Report:")
print(classification_report(y_test, y_pred))
```





Correlation Matrix:

| | Age | YearsAtCompany | MonthlyIncome | Attrition |
|----------------|-----------|----------------|---------------|-----------|
| Age | 1.000000 | 0.311309 | 0.497855 | -0.159205 |
| YearsAtCompany | 0.311309 | 1.000000 | 0.514285 | -0.134392 |
| MonthlyIncome | 0.497855 | 0.514285 | 1.000000 | -0.159840 |
| Attrition | -0.159205 | -0.134392 | -0.159840 | 1.000000 |



Logistic Regression Coefficients (after scaling):

Age: -0.2691
YearsAtCompany: -0.2021
MonthlyIncome: -0.2967

Classification Report:

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.90 | 0.58 | 0.70 | 380 |
| 1 | 0.19 | 0.62 | 0.29 | 61 |
| accuracy | | | 0.58 | 441 |
| macro avg | 0.55 | 0.60 | 0.50 | 441 |
| weighted avg | 0.81 | 0.58 | 0.65 | 441 |

```
# Analyze Attrition by Department and JobRole
# Calculate attrition rate by Department
print("\nAttrition Rate by Department:")
department_attrition = df.groupby('Department')['Attrition'].mean().sort_values(ascending=False) * 100
print(department_attrition)

# Visualize attrition by Department
plt.figure(figsize=(10, 6))
sns.barplot(x='Attrition', y='Department', data=df)
plt.title('Attrition Rate by Department (%)')
plt.xlabel('Attrition Rate (%)')
plt.xlim(0, 0.5) # Since Attrition is 0/1, the mean will be between 0 and 1; adjust for better visualization
plt.show()

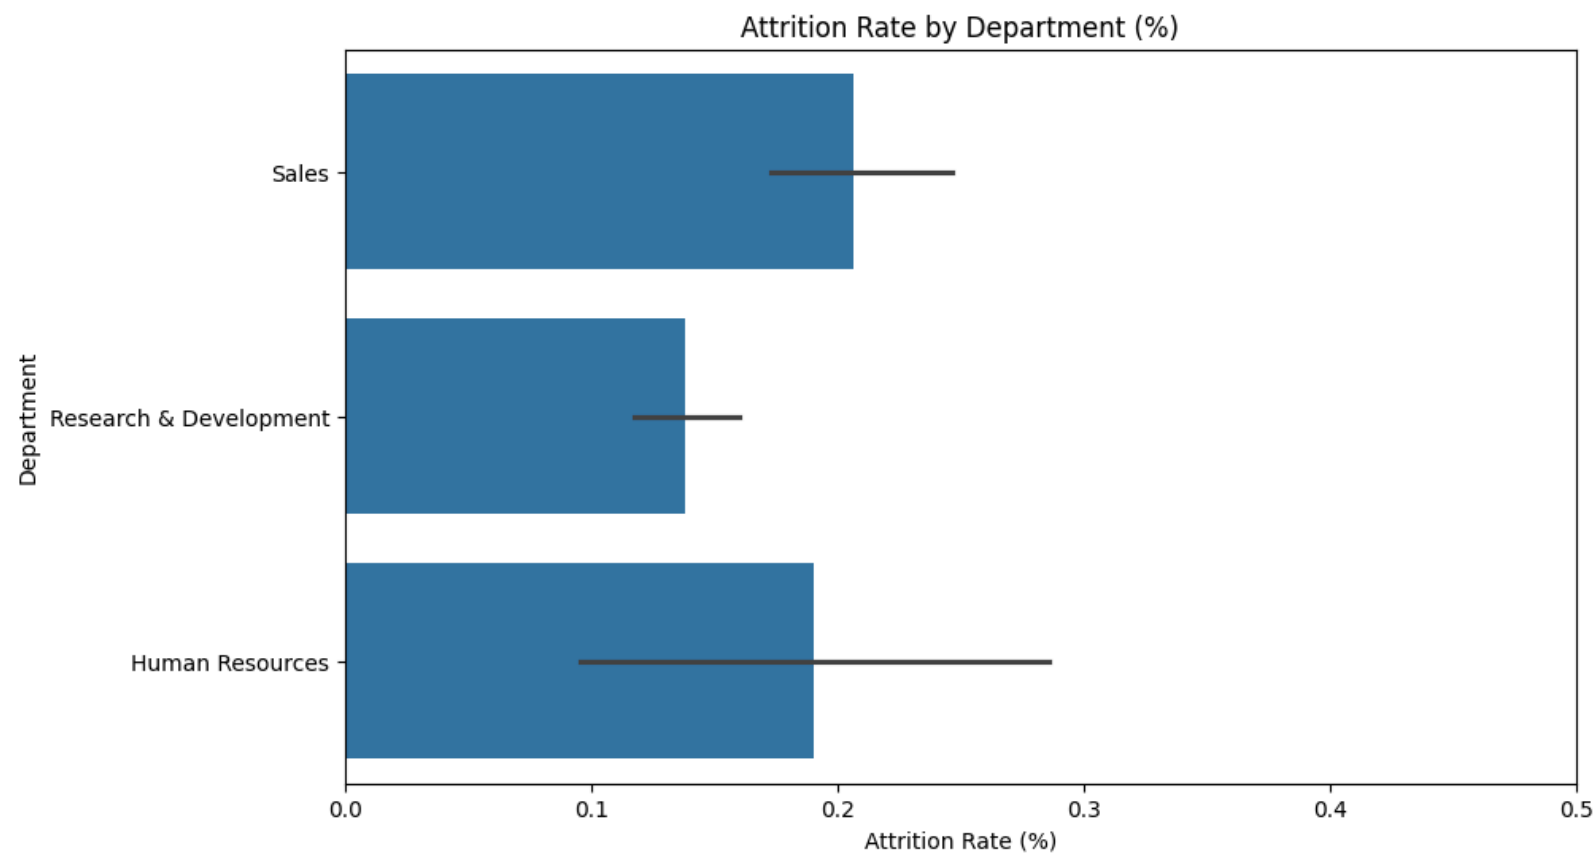
# Calculate attrition rate by JobRole
print("\nAttrition Rate by JobRole:")
jobrole_attrition = df.groupby('JobRole')['Attrition'].mean().sort_values(ascending=False) * 100
print(jobrole_attrition)

# Visualize attrition by JobRole
plt.figure(figsize=(10, 8))
sns.barplot(x='Attrition', y='JobRole', data=df)
plt.title('Attrition Rate by JobRole (%)')
plt.xlabel('Attrition Rate (%)')
plt.xlim(0, 0.5) # Adjust for better visualization
plt.show()
```

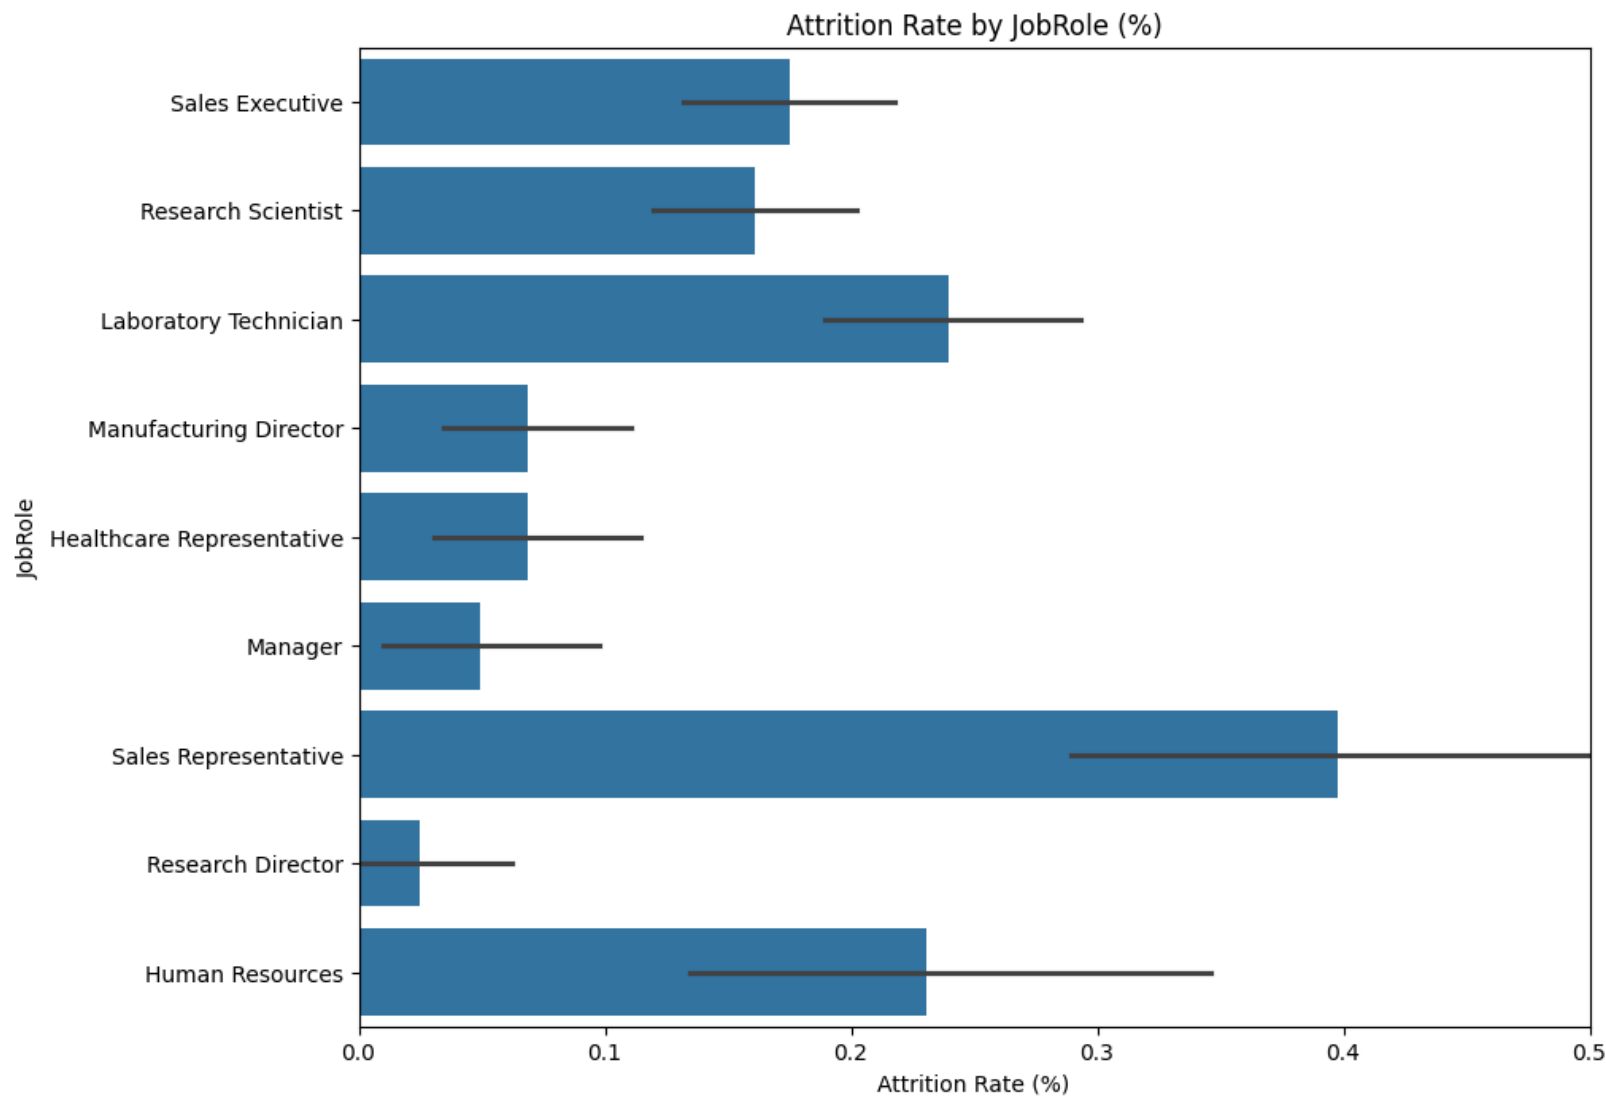




Attrition Rate by Department:
Department
Sales 20.627803
Human Resources 19.047619
Research & Development 13.839750
Name: Attrition, dtype: float64



Attrition Rate by JobRole:
JobRole
Sales Representative 39.759036
Laboratory Technician 23.938224
Human Resources 23.076923
Sales Executive 17.484663
Research Scientist 16.095890
Manufacturing Director 6.896552
Healthcare Representative 6.870229
Manager 4.901961
Research Director 2.500000
Name: Attrition, dtype: float64



```
# Calculate attrition rate by OverTime
print("\nAttrition Rate by OverTime:")
overtime_attrition = df.groupby('OverTime')['Attrition'].mean().sort_values(ascending=False) * 100
print(overtime_attrition)
```

```
# Visualize attrition by OverTime
plt.figure(figsize=(8, 5))
sns.barplot(x='Attrition', y='OverTime', data=df)
plt.title('Attrition Rate by OverTime (%)')
plt.xlabel('Attrition Rate (%)')
plt.xlim(0, 0.5) # Since Attrition is 0/1, the mean will be between 0 and 1
plt.show()
```

```
# Analyze WorkLifeBalance by Attrition
print("\nAverage WorkLifeBalance by Attrition:")
```

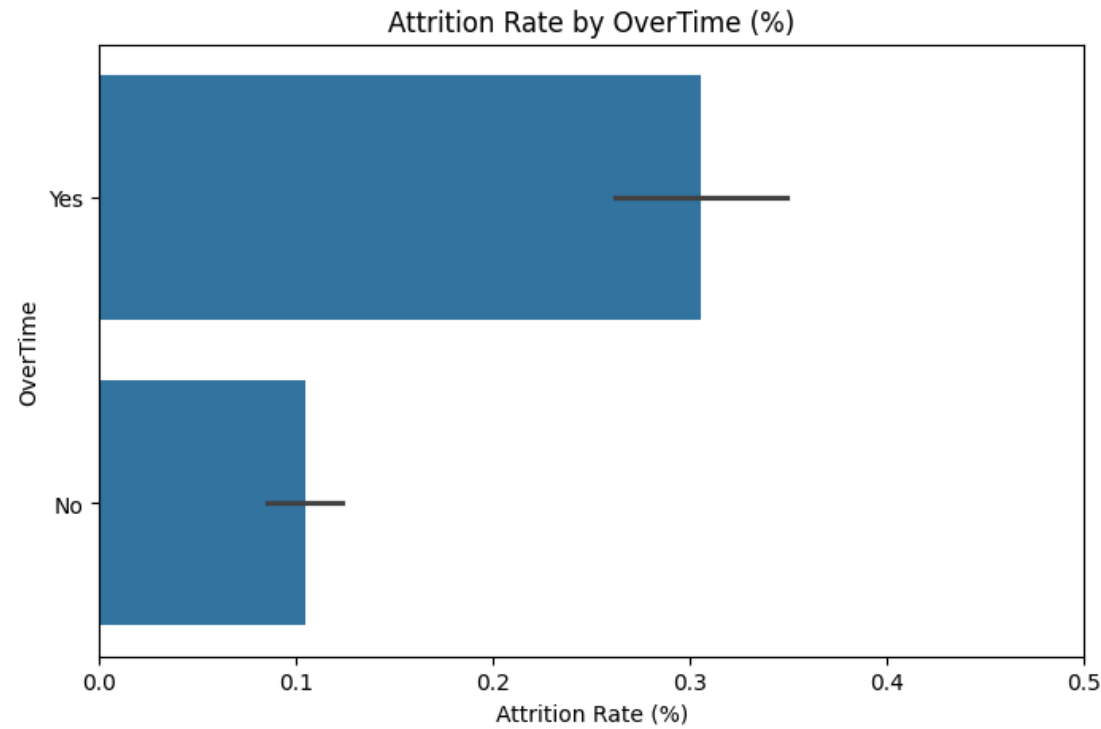


```
print(average_worklifebalance_by_attrition)
worklifebalance_attrition = df.groupby('Attrition')['WorkLifeBalance'].mean()
print(worklifebalance_attrition)
```

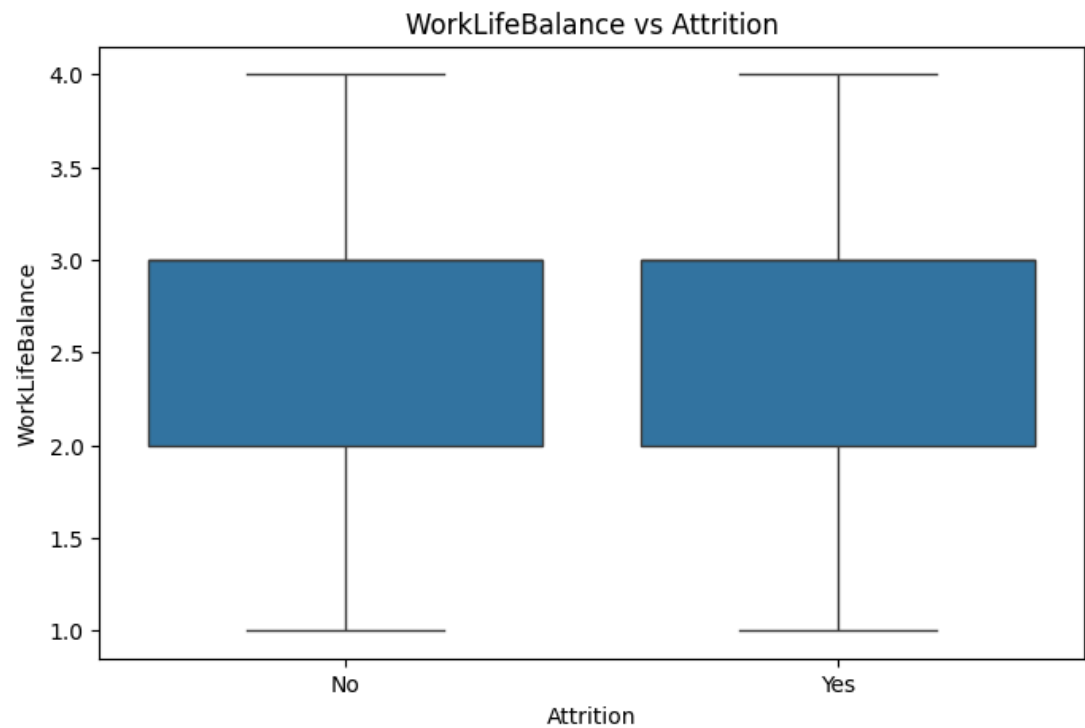
```
# Visualize WorkLifeBalance vs Attrition
plt.figure(figsize=(8, 5))
sns.boxplot(x='Attrition', y='WorkLifeBalance', data=df)
plt.title('WorkLifeBalance vs Attrition')
plt.xticks([0, 1], ['No', 'Yes']) # Label Attrition as No/Yes for clarity
plt.show()
```



Attrition Rate by OverTime:
OverTime
Yes 30.528846
No 10.436433
Name: Attrition, dtype: float64



Average WorkLifeBalance by Attrition:
Attrition
0 2.781022
1 2.658228
Name: WorkLifeBalance, dtype: float64



```
# Add a column for Attrition as a string for Tableau
df['Attrition_Str'] = df['Attrition'].map({1: 'Yes', 0: 'No'})
df.to_csv('processed_employee_data.csv', index=False)
```

